



Deep Learning School

Физтех-Школа Прикладной математики и информатики (ФПМИ) МФТИ

Для быстрого выполнения просмотрите [семинар](#).

▼ Models: Sentence Sentiment Classification

Our goal is to create a model that takes a sentence (just like the ones in our dataset) and produces either 1 (indicating the sentence carries a positive sentiment) or a 0 (indicating the sentence carries a negative sentiment). We can think of it as looking like this:




Under the hood, the model is actually made up of two model.

- DistilBERT processes the sentence and passes along some information it extracted from it on to the next model. DistilBERT is a smaller version of BERT developed and open sourced

by the team at HuggingFace. It's a lighter and faster version of BERT that roughly matches its performance.

- The next model, a basic Logistic Regression model from scikit learn will take in the result of DistilBERT's processing, and classify the sentence as either positive or negative (1 or 0, respectively).

The data we pass between the two models is a vector of size 768. We can think of this of vector as an embedding for the sentence that we can use for classification.



Dataset

The dataset we will use in this example is [SST2](#), which contains sentences from movie reviews, each labeled as either positive (has the value 1) or negative (has the value 0):

sentence	label
a stirring , funny and finally transporting re imagining of beauty and the beast and 1930s horror films	1
apparently reassembled from the cutting room floor of any given daytime soap	0
they presume their audience won't sit still for a sociology lesson	0
this is a visually stunning rumination on love , memory , history and the war between art and commerce	1
jonathan parker 's bartleby should have been the be all end all of the modern office anomie films	1

Installing the transformers library

Let's start by installing the huggingface transformers library so we can load our deep learning NLP model.

```
!pip install transformers
```

[Transformers library doc](#)



HUGGING FACE

On a mission to solve NLP,
one commit at a time.



Star

36,299

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
import torch
import transformers as ppb
import warnings
warnings.filterwarnings('ignore')
```

!nvidia-smi

Mon Nov 22 22:00:37 2021

NVIDIA-SMI		450.119.04		Driver Version: 450.119.04			CUDA Version: 11.0		
-----+-----									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr.	ECC	
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.	MIG M.	
=====+=====									
0	Tesla	P100-PCIE...	Off	00000000:00:04.0	Off				0
N/A	60C	P0	39W / 250W	3511MiB / 16280MiB		0%	Default		N/A
-----+-----									

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage	
	ID	ID					

▼ Importing the dataset

```
df = pd.read_csv(
    'https://github.com/clairett/pytorch-sentiment-classification/raw/master/data/'
    delimiter='\t',
    header=None
)
print(df.shape)
df.head()
```

(6920, 2)

		0	1
0	a stirring , funny and finally transporting re...	1	
1	apparently reassembled from the cutting room f...	0	
2	they presume their audience wo n't sit still f...	0	
3	this is a visually stunning rumination on love...	1	
4	jonathan parker 's bartleby should have been t...	1	

```
np.array(df[1])[1]
```

0

▼ Using BERT for text classification.

Let's now load a pre-trained BERT model.

```
# For DistilBERT, Load pretrained model/tokenizer:
```

```
model_class, tokenizer_class, pretrained_weights = (ppb.DistilBertModel, ppb.Disti
tokenizer = tokenizer_class.from_pretrained(pretrained_weights)
model = model_class.from_pretrained(pretrained_weights)
```

```
# look at the model
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)
model.eval()
```

```
(sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
```

```
print(tokenizer.encode("Hi"))
```

```
[101, 7632, 102]
```

```
from termcolor import colored

colors = ['red', 'green', 'blue', 'yellow']

def model_structure(layer, margin=0, item_color=0):
    for name, next_layer in layer.named_children():

        next = (0 if not list(next_layer.named_children()) else 1)
        print(colored(' ' * margin + name, colors[item_color]) + ':' * next)
        model_structure(next_layer, margin + len(name) + 2, (item_color + 1) % 4)

# model_structure(model)
```

▼ Preparing the dataset

```
from torch.utils.data import Dataset, random_split

class ReviewsDataset(Dataset):
    def __init__(self, reviews, tokenizer, labels):
        self.labels = labels
        # tokenized reviews
        self.tokenized = {i: tokenizer.encode(reviews[i]) for i in range(len(reviews))}

    def __getitem__(self, idx):
        return {"tokenized": self.tokenized[idx], "label": self.labels[idx]}

    def __len__(self):
        return len(self.labels)

dataset = ReviewsDataset(np.array(df[0]), tokenizer, np.array(df[1]))

# DON'T CHANGE, PLEASE
train_size, val_size = int(.8 * len(dataset)), int(.1 * len(dataset))
torch.manual_seed(2)
train_data, valid_data, test_data = random_split(dataset, [train_size, val_size, len(dataset) - train_size - val_size])

print(f"Number of training examples: {len(train_data)}")
print(f"Number of validation examples: {len(valid_data)}")
print(f"Number of testing examples: {len(test_data)}")

Number of training examples: 5536
Number of validation examples: 692
Number of testing examples: 692
```

```
from torch.utils.data import Sampler

class ReviewsSampler(Sampler):
    def __init__(self, subset, batch_size=32):
        self.batch_size = batch_size
```

```

self.subset = subset

self.indices = subset.indices
# tokenized for our data
self.tokenized = []
for i in self.indices:
    self.tokenized.append(subset.dataset.tokenized[i])
self.tokenized = np.array(self.tokenized)
def __iter__(self):

    batch_idx = []
    # index in sorted data
    for index in np.argsort(list(map(len, self.tokenized))):
        batch_idx.append(index)
        if len(batch_idx) == self.batch_size:
            yield batch_idx
            batch_idx = []

    if len(batch_idx) > 0:
        yield batch_idx

def __len__(self):
    return len(self.dataset)

```

```

a = []
for i in train_data.indices:
    a.append(np.array(train_data.dataset.tokenized[i]))
a = np.array(a)
print(a)
print(a.shape)

```

```

[array([ 101, 11552, 2135, 2550, 1998, 22570, 2135, 2864, 1010,
        1996, 2416, 3315, 3616, 6121, 3669, 4371, 3145, 5436,
        5312, 2046, 3371, 2135, 6851, 16278, 1997, 3959, 10359,
        19069, 102])
 array([ 101, 2074, 2178, 12391, 3689, 2008, 2038, 2498, 2183,
        2005, 2009, 2060, 2084, 2049, 18077, 3512, 9140, 1997,
        26471, 10036, 16959, 2015, 102])
 array([ 101, 1037, 10973, 17197, 4038, 2895, 17037, 2008, 1005,
        2222, 2404, 13606, 2006, 2115, 3108, 102])
... array([ 101, 12391, 10874, 18015, 102])
 array([ 101, 1037, 2143, 1997, 8680, 4094, 2007, 2019, 10305,
        3110, 1010, 1037, 12312, 1997, 1996, 11139, 1998, 12482,
        1997, 28956, 102])
 array([ 101, 2025, 2012, 2035, 3154, 2054, 2009, 1005, 1055, 2667, 2000,
        2360, 1998, 2130, 2065, 2009, 2020, 1045, 4797, 2009, 2052, 2022,
        2035, 2008, 5875, 102])]
(5536,)

```

```

from torch.utils.data import DataLoader

```

```

def get_padded(values):
    max_len = 80
    # for value in values:
    #     if len(value) > max_len:

```

```

#         max_len = len(value)

padded = np.array([value + [0]*(max_len-len(value)) for value in values])

return padded

def collate_fn(batch):

    inputs = []
    labels = []
    for elem in batch:
        inputs.append(elem['tokenized'])
        labels.append(elem['label'])

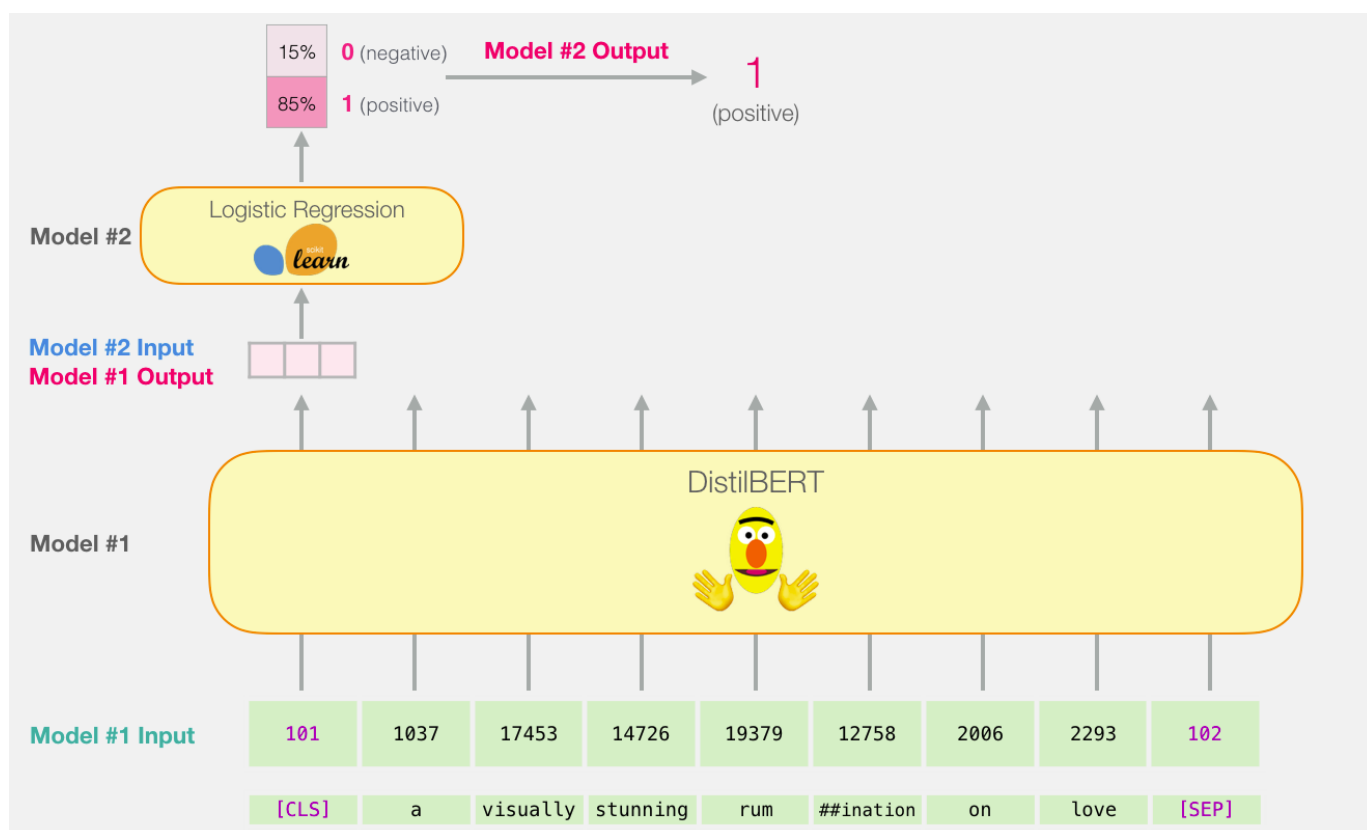
    inputs = get_padded(inputs) # padded inputs
    attention_mask = np.where(inputs != 0, 1, 0)

    return {"inputs": torch.tensor(inputs), "labels": torch.FloatTensor(labels), '

train_loader = DataLoader(train_data, batch_sampler=ReviewsSampler(train_data), co
valid_loader = DataLoader(valid_data, batch_sampler=ReviewsSampler(valid_data), co
test_loader = DataLoader(test_data, batch_sampler=ReviewsSampler(test_data), colla

```

▼ Baseline



```
from tqdm.notebook import tqdm
```

```
def get_xy(loader, model):
```



```

features = []
labels = []

with torch.no_grad():
    for batch in tqdm(loader):

        attn_mask = batch["attention_mask"].to(device)
        inputs = batch["inputs"].to(device)

        with torch.no_grad():
            last_hidden_states = model(inputs, attention_mask=attn_mask)

        features.append(last_hidden_states[0].cpu())
        labels.append(batch["labels"])

features = torch.cat([elem[:, 0, :] for elem in features], dim=0).numpy()
labels = torch.cat(labels, dim=0).numpy()

return features, labels

train_features, train_labels = get_xy(train_loader, model)
valid_features, valid_labels = get_xy(valid_loader, model)
test_features, test_labels = get_xy(test_loader, model)

0it [00:00, ?it/s]
0it [00:00, ?it/s]
0it [00:00, ?it/s]

lr_clf = LogisticRegression(C = 1.2, solver='liblinear')
lr_clf.fit(train_features, train_labels.astype(int))
lr_clf.score(test_features, test_labels)

0.8208092485549133

```

▼ Fine-Tuning BERT

Define the model

```

from torch import nn

class BertClassifier(nn.Module):
    def __init__(self, pretrained_model, hid_dim, dropout=0.1):
        super().__init__()

        self.bert = pretrained_model
        self.dropout = nn.Dropout(p=dropout)
        self.relu = nn.ReLU()
        self.hid_dim = hid_dim

        self.fc = nn.Linear(hid_dim * 80, 1)

        self.sigm = nn.Sigmoid()

```

```

def forward(self, inputs, attention_mask):

    predictions = self.bert(inputs, attention_mask=attention_mask).last_hidden_
    predictions = self.relu(predictions)
    predictions = predictions.reshape(-1, self.hid_dim*80)
    probs = self.fc(predictions)
    probs = self.sigm(probs)

    # proba = [batch_size, ] - probability to be positive
    return probs

import torch.optim as optim

# DON'T CHANGE
model = model_class.from_pretrained(pretrained_weights).to(device)
bert_clf = BertClassifier(model, 768).to(device)
# you can change
optimizer = optim.Adam(bert_clf.parameters(), lr=2e-5)
criterion = nn.BCELoss()

def train(model, iterator, optimizer, criterion, clip, train_history=None, valid_h:
    model.train()

    epoch_loss = 0
    history = []
    for i, batch in enumerate(iterator):

        # don't forget about .to(device)
        inputs = batch["inputs"].to(device)
        mask = batch["attention_mask"].to(device)
        labels = batch["labels"].to(device)

        optimizer.zero_grad()

        output = model(inputs, mask).reshape(-1)

        loss = criterion(output, labels)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), clip)
        optimizer.step()

        epoch_loss += loss.item()

        history.append(loss.cpu().data.numpy())
        if (i+1)%10==0:
            fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(12, 8))

            clear_output(True)
            ax[0].plot(history, label='train loss')
            ax[0].set_xlabel('Batch')
            ax[0].set_title('Train loss')

```

```

        if train_history is not None:
            ax[1].plot(train_history, label='general train history')
            ax[1].set_xlabel('Epoch')
        if valid_history is not None:
            ax[1].plot(valid_history, label='general valid history')
        plt.legend()

    plt.show()

    return epoch_loss / (i + 1)

def evaluate(model, iterator, criterion):

    model.eval()

    epoch_loss = 0

    history = []

    with torch.no_grad():

        for i, batch in enumerate(iterator):

            inputs = batch["inputs"].to(device)
            mask = batch["attention_mask"].to(device)
            labels = batch["labels"].to(device)

            output = model(inputs, mask).reshape(-1)

            loss = criterion(output, labels)

            epoch_loss += loss.item()

    return epoch_loss / (i + 1)

def epoch_time(start_time, end_time):
    elapsed_time = end_time - start_time
    elapsed_mins = int(elapsed_time / 60)
    elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
    return elapsed_mins, elapsed_secs

import time
import math
import matplotlib
matplotlib.rcParams.update({'figure.figsize': (16, 12), 'font.size': 14})
import matplotlib.pyplot as plt
%matplotlib inline
from IPython.display import clear_output

train_history = []
valid_history = []

```

```
N_EPOCHS = 3
CLIP = 1

best_valid_loss = float('inf')

for epoch in range(N_EPOCHS):

    start_time = time.time()

    train_loss = train(bert_clf, train_loader, optimizer, criterion, CLIP, train_h:
    valid_loss = evaluate(bert_clf, valid_loader, criterion)

    end_time = time.time()

    epoch_mins, epoch_secs = epoch_time(start_time, end_time)

    if valid_loss < best_valid_loss:
        best_valid_loss = valid_loss
        torch.save(bert_clf.state_dict(), 'best-val-model.pt')

    train_history.append(train_loss)
    valid_history.append(valid_loss)
    print(f'Epoch: {epoch+1:02} | Time: {epoch_mins}m {epoch_secs}s')
    print(f'\tTrain Loss: {train_loss:.3f} | Train PPL: {math.exp(train_loss):7.3f}')
    print(f'\tVal. Loss: {valid_loss:.3f} | Val. PPL: {math.exp(valid_loss):7.3f}')
```

```

Train loss
a = torch.tensor([1, 2, 3, 4, 5])
print(np.array((a>3).to(torch.int)))

[0 0 0 1 1]
0.5
best_model = BertClassifier(model, 768).to(device)
best_model.load_state_dict(torch.load('best-val-model.pt'))

pred_labels = []
true_labels = []

best_model.eval()
with torch.no_grad():
    for i, batch in tqdm(enumerate(test_loader)):
        inputs = batch["inputs"].to(device)
        mask = batch["attention_mask"].to(device)
        labels = batch["labels"].to(device)

        true_labels.append(labels.cpu().numpy())

        output = best_model(inputs, mask).reshape(-1)
        pred_labels.append(np.array((output>0.5).to(torch.int).cpu()))

0it [00:00, ?it/s]
val loss: 0.170 | val f1: 0.877

from sklearn.metrics import accuracy_score

true_labels = np.concatenate(true_labels, axis=0)
pred_labels = np.concatenate(pred_labels, axis=0)
accuracy_score(true_labels, pred_labels)

0.865606936416185

assert accuracy_score(true_labels, pred_labels) >= 0.86

```

▼ Finetuned model from **HUGGING FACE**

[BertForSequenceClassification](#)

```

from transformers import AutoTokenizer, AutoModelForSequenceClassification

# we have the same tokenizer
# new_tokenizer = AutoTokenizer.from_pretrained("distilbert-base-uncased-finetuned-
new_model = AutoModelForSequenceClassification.from_pretrained("distilbert-base-uncased-finetuned-

Downloading: 0%|          | 0.00/629 [00:00<?, ?B/s]
Downloading: 0%|          | 0.00/268M [00:00<?, ?B/s]

pred_labels = []

```

```

true_labels = []

new_model.eval()
with torch.no_grad():
    for i, batch in tqdm(enumerate(test_loader)):
        inputs = batch["inputs"].to(device)
        mask = batch["attention_mask"].to(device)
        labels = batch["labels"].to(device)

        true_labels.append(labels.cpu().numpy())

    output = new_model(inputs, mask).logits
    pred_labels.append(np.array(torch.argmax(output, dim=-1).cpu()))

true_labels = np.concatenate(true_labels, axis=0)
pred_labels = np.concatenate(pred_labels, axis=0)
accuracy_score(true_labels, pred_labels)

0it [00:00, ?it/s]
0.9841040462427746

```

```
model_structure(new_model)
```



```

        dropout
        lin1
        lin2

    output_layer_norm
2:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
3:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
4:
    attention:
        dropout

```

```

        q_lin
        k_lin
        v_lin
        out_lin

    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
5:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin

    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm

pre_classifier
classifier
dropout

```

Напишите вывод о своих результатах. В выводы включите ваши гиперпараметры.

Качество с помощью Fine-Tuning должно достигать 0.86.

На классификации с логистической регрессии точность вышла 0.82
 На Fine-Tuning - 0.86. Параметры: max_len=80
 На FineTuned - 0.984

